

CNN-LSTM Parallel Moneymaker

```
In [1]: import torch
import torch.optim
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader

import time
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: X_train, X_val, y_train, y_val, train_mean, val_mean, train_std, val_std
= torch.load("assets/all.pt")
```

```
In [3]: print("X_train shape: \t\t", X_train.shape)
print("X_val shape: \t\t", X_val.shape)
print("y_train shape: \t\t", y_train.shape)
print("y_val shape: \t\t", y_val.shape)
```

```
X_train shape:          torch.Size([2542795, 122, 3])
X_val shape:            torch.Size([282732, 122, 3])
y_train shape:          torch.Size([2542795, 1])
y_val shape:            torch.Size([282732, 1])
```

```
In [4]: # optimize for GPU if exists
if torch.cuda.is_available():
    X_train = X_train.cuda()
    y_train = y_train.cuda()
    X_val = X_val.cuda()
    y_val = y_val.cuda()
```

```
In [5]: N, S, D = X_train.shape
```

Training a CNN

```
In [6]: class Dataset(torch.utils.data.Dataset):
    def __init__(self, X, y):
        self.X = X
        self.y = y

    def __len__(self):
        return len(self.X)

    def __getitem__(self, index):
        return self.X[index], self.y[index]
```

```
In [7]: batch_size = 32
dataset = Dataset(X_train, y_train)
loader = DataLoader(dataset, batch_size, shuffle=True)
```

```
In [8]: class CNNClassifier(nn.Module):
    def __init__(self, input_dim, hidden_dim, output_dim):
        super(CNNClassifier, self).__init__()
        self.lstm = nn.LSTM(input_dim, hidden_dim, batch_first=True)
        self.drop = nn.Dropout()
        self.conv1 = nn.Conv1d(input_dim, hidden_dim, 7)
        self.pool = nn.MaxPool1d(3)
        self.flattened_dim = int((S - 7) / 3) * hidden_dim

        self.fc = nn.Linear(self.flattened_dim + hidden_dim, output_dim)

    def forward(self, x, h=None):
        # LSTM
        if type(h) == type(None):
            x1, hn = self.lstm(x)
        else:
            x1, hn = self.lstm(x, h.detach())
        x1 = x1[:, -1, :]
        x1 = self.drop(x1)

        # Conv
        x2 = torch.transpose(x, 1, 2)
        x2 = self.pool(F.relu(self.conv1(x2)))
        x2 = x2.view(-1, self.flattened_dim)

        x = torch.cat([x1, x2], dim=1)
        out = self.fc(x)
        return out
```

```
In [9]: input_dim = 3
hidden_dim = 32
output_dim = 1
```

```
In [10]: model = CNNClassifier(input_dim, hidden_dim, output_dim)
if torch.cuda.is_available():
    model = model.to("cuda")

criterion = nn.MSELoss()
optimizer = torch.optim.Adam(model.parameters())
```

```
In [11]: train_accs = []
val_accs = []
train_losses = []
val_losses = []
epoch = 0
```

Train the model

```
In [12]: def pred_val(X_val, model):  
    val_batch_size = 1000  
    val_set_size = X_val.shape[0]  
    preds = []  
    with torch.no_grad():  
        for i in range(0, val_set_size, val_batch_size):  
            start = i  
            end = min(i+val_batch_size, val_set_size)  
            preds.append(model(X_val[start:end]))  
    pred = torch.cat(preds, dim=0)  
    return pred
```

```

In [13]: t0 = time.time()
num_epochs = 3
for ep in range(num_epochs):
    tstart = time.time()
    for i, data in enumerate(loader):
        model.train()
        print("{} / {}".format(i, int(X_train.shape[0] / batch_size)), end
        = '\r')
        optimizer.zero_grad()
        outputs = model(data[0])
        loss = criterion(outputs, data[1])
        loss.backward()
        optimizer.step()

    if i % 2500 == 0:
        with torch.no_grad():
            model.eval()
            train_losses.append(loss.item())
            pXval = pred_val(X_val, model)
            vloss = criterion(pXval, y_val)
            val_losses.append(vloss.item())
            torch.save({
                'epoch': epoch,
                'model_state_dict': model.state_dict(),
                'optimizer_state_dict': optimizer.state_dict(),
                'loss': loss,
            }, 'assets/partial_model.pt')
            print("training loss: {:<3.3f} \t val loss: {:<3.3f}".fo
rmat(loss, vloss))

        with torch.no_grad():
            model.eval()
            pXval = pred_val(X_val, model)
            vloss = criterion(pXval, y_val)
            val_losses.append(vloss.item())
            epoch += 1
            tend = time.time()
            print('epoch: {:<3d} \t time: {:<3.2f} \t val loss: {:<3.3f}'.fo
rmat(epoch,
            tend - tstart, vloss.item()))
time_total = time.time() - t0
print('Total time: {:<4.3f}, average time per epoch: {:<4.3f}'.format(time
_total, time_total / num_epochs))

```


| | | | |
|----------------|--------------|-----------|-------|
| training loss: | 2.773 | val loss: | 1.913 |
| training loss: | 0.227 | val loss: | 0.161 |
| training loss: | 0.167 | val loss: | 0.151 |
| training loss: | 0.110 | val loss: | 0.131 |
| training loss: | 0.126 | val loss: | 0.141 |
| training loss: | 0.133 | val loss: | 0.126 |
| training loss: | 0.080 | val loss: | 0.129 |
| training loss: | 0.364 | val loss: | 0.122 |
| training loss: | 0.165 | val loss: | 0.130 |
| training loss: | 0.089 | val loss: | 0.113 |
| training loss: | 0.075 | val loss: | 0.117 |
| training loss: | 0.156 | val loss: | 0.109 |
| training loss: | 0.070 | val loss: | 0.108 |
| training loss: | 0.080 | val loss: | 0.103 |
| training loss: | 0.050 | val loss: | 0.106 |
| training loss: | 0.084 | val loss: | 0.104 |
| training loss: | 0.109 | val loss: | 0.105 |
| training loss: | 0.107 | val loss: | 0.101 |
| training loss: | 0.100 | val loss: | 0.114 |
| training loss: | 0.133 | val loss: | 0.119 |
| training loss: | 0.114 | val loss: | 0.115 |
| training loss: | 0.087 | val loss: | 0.100 |
| training loss: | 0.082 | val loss: | 0.110 |
| training loss: | 0.351 | val loss: | 0.118 |
| training loss: | 0.046 | val loss: | 0.095 |
| training loss: | 0.056 | val loss: | 0.094 |
| training loss: | 0.113 | val loss: | 0.107 |
| training loss: | 0.091 | val loss: | 0.096 |
| training loss: | 0.122 | val loss: | 0.101 |
| training loss: | 0.147 | val loss: | 0.088 |
| training loss: | 0.077 | val loss: | 0.104 |
| training loss: | 0.050 | val loss: | 0.089 |
| epoch: 1 | time: 576.77 | val loss: | 0.089 |
| training loss: | 0.244 | val loss: | 0.100 |
| training loss: | 0.076 | val loss: | 0.089 |
| training loss: | 0.043 | val loss: | 0.093 |
| training loss: | 0.110 | val loss: | 0.088 |
| training loss: | 0.046 | val loss: | 0.104 |
| training loss: | 0.101 | val loss: | 0.101 |
| training loss: | 0.086 | val loss: | 0.085 |
| training loss: | 0.121 | val loss: | 0.099 |
| training loss: | 0.051 | val loss: | 0.093 |
| training loss: | 0.061 | val loss: | 0.084 |
| training loss: | 0.106 | val loss: | 0.106 |
| training loss: | 0.060 | val loss: | 0.097 |
| training loss: | 0.142 | val loss: | 0.083 |
| training loss: | 0.101 | val loss: | 0.088 |
| training loss: | 0.098 | val loss: | 0.088 |
| training loss: | 0.104 | val loss: | 0.085 |
| training loss: | 0.117 | val loss: | 0.083 |
| training loss: | 0.072 | val loss: | 0.089 |
| training loss: | 0.055 | val loss: | 0.094 |
| training loss: | 0.082 | val loss: | 0.088 |
| training loss: | 0.074 | val loss: | 0.080 |
| training loss: | 0.081 | val loss: | 0.091 |
| training loss: | 0.077 | val loss: | 0.081 |
| training loss: | 0.047 | val loss: | 0.085 |

```
training loss: 0.055    val loss: 0.088
training loss: 0.144    val loss: 0.093
training loss: 0.048    val loss: 0.080
training loss: 0.044    val loss: 0.101
training loss: 0.067    val loss: 0.088
training loss: 0.120    val loss: 0.091
training loss: 0.064    val loss: 0.080
training loss: 0.070    val loss: 0.098
epoch: 2              time: 576.35    val loss: 0.080
training loss: 0.072    val loss: 0.078
training loss: 0.053    val loss: 0.084
training loss: 0.106    val loss: 0.104
training loss: 0.064    val loss: 0.098
training loss: 0.064    val loss: 0.079
training loss: 0.054    val loss: 0.082
training loss: 0.036    val loss: 0.079
training loss: 0.058    val loss: 0.076
training loss: 0.053    val loss: 0.083
training loss: 0.060    val loss: 0.077
training loss: 0.098    val loss: 0.079
training loss: 0.076    val loss: 0.079
training loss: 0.114    val loss: 0.100
training loss: 0.052    val loss: 0.074
training loss: 0.067    val loss: 0.087
training loss: 0.063    val loss: 0.083
training loss: 0.088    val loss: 0.091
training loss: 0.076    val loss: 0.078
training loss: 0.117    val loss: 0.076
training loss: 0.052    val loss: 0.078
training loss: 0.074    val loss: 0.073
training loss: 0.079    val loss: 0.084
training loss: 0.066    val loss: 0.076
training loss: 0.060    val loss: 0.076
training loss: 0.074    val loss: 0.089
training loss: 0.069    val loss: 0.076
training loss: 0.079    val loss: 0.087
training loss: 0.052    val loss: 0.086
training loss: 0.064    val loss: 0.077
training loss: 0.045    val loss: 0.076
training loss: 0.056    val loss: 0.079
training loss: 0.049    val loss: 0.075
epoch: 3              time: 574.98    val loss: 0.073
Total time: 1728.102, average time per epoch: 576.034
```

Training loss vs. validation loss

```
In [19]: t_losses = [i for i in train_losses if i < 4000]
plt.plot(t_losses)
plt.plot(val_losses)
plt.title('loss history')
plt.xlabel('2500 iterations')
plt.ylabel('loss')
plt.yscale('log')
plt.legend(['train', 'val'])

plt.show()
```



Evaluate the model

```
In [16]: X_train = X_train.cuda()
y_train = y_train.cuda()
X_val = X_val.cuda()
y_val = y_val.cuda()
```

```
In [17]: model.eval()

pred = pred_val(X_val, model)
val_loss = criterion(pred, y_val).item()

print("Final model evaluation: ", val_loss)

Predicting 282000 - 282732...
Final model evaluation: 0.0734802633523941
```

One-step lag predictor

The one-step lag predictor simply outputs the last timestep in the input sequence. Our model should outperform the one-step lag predictor.

```
In [20]: def one_step_lag_predictor(X):  
         return X[:, -1, 0].unsqueeze(1)  
  
p_val_naive = one_step_lag_predictor(X_val.cpu())  
loss_naive = criterion(p_val_naive, y_val.cpu())  
  
print("Loss from 1-step lag predictor:\t{}\nLoss from our model:\t\t{}".  
      format(loss_naive, val_loss))
```

```
Loss from 1-step lag predictor: 0.14193207025527954  
Loss from our model:          0.0734802633523941
```

Standard deviation difference

```
In [21]: # switch back to cpu for plotting  
X_train = X_train.cpu()  
y_train = y_train.cpu()  
X_val = X_val.cpu()  
y_val = y_val.cpu()  
pred = pred.cpu()  
  
# backprop components no longer needed  
X_train = X_train.detach()  
y_train = y_train.detach()  
X_val = X_val.detach()  
y_val = y_val.detach()  
pred = pred.detach()
```




```
In [22]: f1 = plt.figure()

ax1 = f1.add_subplot()
ax1.plot((pred - y_val)[500:522])
ax1.set_title('price std diff')
ax1.set(xlabel='window', ylabel='std')

plt.show()

# plt.plot((pred[:,3] - y_val.cpu()[:,3]).detach())
# plt.title('std difference')
# plt.plot([1, 2, 3])
```



```
In [24]: # denormalize the data
pred_abs = pred * val_std[:,0].unsqueeze(1) + val_mean[:,0].unsqueeze(1)
y_val_abs = y_val.cpu() * val_std[:,0].unsqueeze(1) + val_mean[:,0].unsq
ueeze(1)
```

```
In [34]: start_window = 500

fig, (ax1, ax2) = plt.subplots(2, figsize=(15, 10))
ax1.plot((pred_abs - y_val_abs)[start_window:start_window + 60])
l1, = ax2.plot(pred_abs[start_window:start_window + 60])
l1.set_label("predicted price")
l2, = ax2.plot(y_val_abs[start_window:start_window + 60])
l2.set_label("actual price")

plt.legend()
plt.show()
```

