

CNN Moneymaker

```
In [1]: import torch
import torch.optim
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader

import time
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: X_train, X_val, y_train, y_val, train_mean, val_mean, train_std, val_std
= torch.load("assets/all.pt")
```

```
In [3]: print("X_train shape: \t\t", X_train.shape)
print("X_val shape: \t\t", X_val.shape)
print("y_train shape: \t\t", y_train.shape)
print("y_val shape: \t\t", y_val.shape)
```

```
X_train shape:          torch.Size([2542795, 122, 3])
X_val shape:            torch.Size([282732, 122, 3])
y_train shape:          torch.Size([2542795, 1])
y_val shape:            torch.Size([282732, 1])
```

```
In [4]: # optimize for GPU if exists
if torch.cuda.is_available():
    X_train = X_train.cuda()
    y_train = y_train.cuda()
    X_val = X_val.cuda()
    y_val = y_val.cuda()
```

```
In [5]: N, S, D = X_train.shape
```

Training a CNN

```
In [6]: class Dataset(torch.utils.data.Dataset):
    def __init__(self, X, y):
        self.X = X
        self.y = y

    def __len__(self):
        return len(self.X)

    def __getitem__(self, index):
        return self.X[index], self.y[index]
```

```
In [7]: batch_size = 32
dataset = Dataset(X_train, y_train)
loader = DataLoader(dataset, batch_size, shuffle=True)
```

```
In [8]: class CNNClassifier(nn.Module):
    def __init__(self, input_dim, hidden_dim, output_dim):
        super(CNNClassifier, self).__init__()
        self.conv1 = nn.Conv1d(input_dim, hidden_dim, 7)
        self.pool = nn.MaxPool1d(3)
        self.flattened_dim = int((S - 7 + 1) / 3) * hidden_dim

        self.fc = nn.Linear(self.flattened_dim, output_dim)

    def forward(self, x, h=None):

        # Conv
        x2 = torch.transpose(x, 1, 2)
        x2 = self.pool(F.relu(self.conv1(x2)))
        x2 = x2.view(-1, self.flattened_dim)

        out = self.fc(x2)
        return out
```

```
In [9]: input_dim = 3
hidden_dim = 32
output_dim = 1
```

```
In [10]: model = CNNClassifier(input_dim, hidden_dim, output_dim)
if torch.cuda.is_available():
    model = model.to("cuda")

criterion = nn.MSELoss()
optimizer = torch.optim.Adam(model.parameters())
```

```
In [11]: train_accs = []
val_accs = []
train_losses = []
val_losses = []
epoch = 0
```

Train the model

```
In [12]: def pred_val(X_val, model):  
    val_batch_size = 1000  
    val_set_size = X_val.shape[0]  
    preds = []  
    with torch.no_grad():  
        for i in range(0, val_set_size, val_batch_size):  
            start = i  
            end = min(i+val_batch_size, val_set_size)  
            preds.append(model(X_val[start:end]))  
    pred = torch.cat(preds, dim=0)  
    return pred
```

```

In [13]: t0 = time.time()
num_epochs = 3
for ep in range(num_epochs):
    tstart = time.time()
    for i, data in enumerate(loader):
        model.train()
        print("{} / {}".format(i, int(X_train.shape[0] / batch_size)), end
        ='\r')
        optimizer.zero_grad()
        outputs = model(data[0])
        loss = criterion(outputs, data[1])
        loss.backward()
        optimizer.step()

    if i % 2500 == 0:
        with torch.no_grad():
            model.eval()
            train_losses.append(loss.item())
            pXval = pred_val(X_val, model)
            vloss = criterion(pXval, y_val)
            val_losses.append(vloss.item())
            torch.save({
                'epoch': epoch,
                'model_state_dict': model.state_dict(),
                'optimizer_state_dict': optimizer.state_dict(),
                'loss': loss,
            }, 'assets/partial_model.pt')
            print("training loss: {:<3.3f} \t val loss: {:<3.3f}".fo
rmat(loss, vloss))

        with torch.no_grad():
            model.eval()
            pXval = pred_val(X_val, model)
            vloss = criterion(pXval, y_val)
            val_losses.append(vloss.item())
            epoch += 1
            tend = time.time()
            print('epoch: {:<3d} \t time: {:<3.2f} \t val loss: {:<3.3f}'.fo
rmat(epoch,
            tend - tstart, vloss.item()))
time_total = time.time() - t0
print('Total time: {:<4.3f}, average time per epoch: {:<4.3f}'.format(time
_total, time_total / num_epochs))

```


training loss:	1.675	val loss:	1.730
training loss:	0.119	val loss:	0.246
training loss:	0.129	val loss:	0.190
training loss:	0.140	val loss:	0.171
training loss:	0.073	val loss:	0.144
training loss:	0.078	val loss:	0.154
training loss:	0.176	val loss:	0.133
training loss:	0.083	val loss:	0.119
training loss:	0.079	val loss:	0.123
training loss:	0.061	val loss:	0.125
training loss:	0.253	val loss:	0.117
training loss:	0.092	val loss:	0.108
training loss:	0.103	val loss:	0.113
training loss:	0.059	val loss:	0.113
training loss:	0.067	val loss:	0.104
training loss:	0.167	val loss:	0.105
training loss:	0.091	val loss:	0.119
training loss:	0.074	val loss:	0.096
training loss:	0.094	val loss:	0.105
training loss:	0.045	val loss:	0.106
training loss:	0.064	val loss:	0.103
training loss:	0.073	val loss:	0.096
training loss:	0.100	val loss:	0.103
training loss:	0.097	val loss:	0.153
training loss:	0.075	val loss:	0.093
training loss:	0.223	val loss:	0.095
training loss:	0.052	val loss:	0.099
training loss:	0.113	val loss:	0.104
training loss:	0.177	val loss:	0.099
training loss:	0.111	val loss:	0.109
training loss:	0.115	val loss:	0.094
training loss:	0.066	val loss:	0.097
epoch: 1	time: 169.09	val loss:	0.113
training loss:	0.935	val loss:	0.125
training loss:	0.053	val loss:	0.094
training loss:	0.067	val loss:	0.095
training loss:	0.066	val loss:	0.093
training loss:	0.102	val loss:	0.092
training loss:	0.083	val loss:	0.091
training loss:	0.094	val loss:	0.098
training loss:	0.050	val loss:	0.103
training loss:	0.060	val loss:	0.098
training loss:	0.160	val loss:	0.098
training loss:	0.089	val loss:	0.090
training loss:	0.100	val loss:	0.107
training loss:	0.089	val loss:	0.092
training loss:	0.051	val loss:	0.089
training loss:	0.080	val loss:	0.090
training loss:	0.061	val loss:	0.092
training loss:	0.129	val loss:	0.090
training loss:	0.063	val loss:	0.098
training loss:	0.074	val loss:	0.086
training loss:	0.108	val loss:	0.091
training loss:	0.122	val loss:	0.089
training loss:	0.097	val loss:	0.125
training loss:	0.054	val loss:	0.098
training loss:	0.076	val loss:	0.102

```
training loss: 0.063    val loss: 0.106
training loss: 0.083    val loss: 0.088
training loss: 0.083    val loss: 0.089
training loss: 0.108    val loss: 0.087
training loss: 0.093    val loss: 0.086
training loss: 0.062    val loss: 0.093
training loss: 0.066    val loss: 0.087
training loss: 0.040    val loss: 0.085
epoch: 2               time: 169.57    val loss: 0.107
training loss: 0.101    val loss: 0.089
training loss: 0.126    val loss: 0.098
training loss: 0.096    val loss: 0.088
training loss: 0.090    val loss: 0.106
training loss: 0.072    val loss: 0.084
training loss: 0.066    val loss: 0.085
training loss: 0.080    val loss: 0.122
training loss: 0.095    val loss: 0.084
training loss: 0.091    val loss: 0.104
training loss: 0.077    val loss: 0.084
training loss: 0.056    val loss: 0.107
training loss: 0.146    val loss: 0.088
training loss: 0.178    val loss: 0.082
training loss: 0.049    val loss: 0.103
training loss: 0.126    val loss: 0.092
training loss: 0.071    val loss: 0.084
training loss: 0.123    val loss: 0.118
training loss: 0.135    val loss: 0.082
training loss: 0.052    val loss: 0.087
training loss: 0.044    val loss: 0.085
training loss: 0.073    val loss: 0.090
training loss: 0.031    val loss: 0.086
training loss: 0.093    val loss: 0.087
training loss: 0.040    val loss: 0.092
training loss: 0.034    val loss: 0.081
training loss: 0.060    val loss: 0.085
training loss: 0.051    val loss: 0.086
training loss: 0.042    val loss: 0.081
training loss: 0.081    val loss: 0.095
training loss: 0.113    val loss: 0.086
training loss: 0.066    val loss: 0.092
training loss: 0.061    val loss: 0.081
epoch: 3               time: 166.90    val loss: 0.088
Total time: 505.562, average time per epoch: 168.521
```

Training loss vs. validation loss

```
In [14]: t_losses = [i for i in train_losses if i < 4000]
plt.plot(t_losses)
plt.plot(val_losses)
plt.title('loss history')
plt.xlabel('2500 iterations')
plt.ylabel('loss')
plt.yscale('log')
plt.legend(['train', 'val'])

plt.show()
```



Evaluate the model

```
In [15]: X_train = X_train.cuda()
y_train = y_train.cuda()
X_val = X_val.cuda()
y_val = y_val.cuda()
```

```
In [16]: model.eval()

pred = pred_val(X_val, model)
val_loss = criterion(pred, y_val).item()

print("Final model evaluation: ", val_loss)

Final model evaluation: 0.08793757110834122
```

One-step lag predictor

The one-step lag predictor simply outputs the last timestep in the input sequence. Our model should outperform the one-step lag predictor.

```
In [25]: def one_step_lag_predictor(X):  
         return X[:, -1, 0].unsqueeze(1)  
  
p_val_naive = one_step_lag_predictor(X_val.cpu())  
loss_naive = criterion(p_val_naive, y_val.cpu())  
  
print("Loss from 1-step lag predictor:\t{}\nLoss from our model:\t\t{}".  
      format(loss_naive, val_loss))
```

```
Loss from 1-step lag predictor: 0.14193207025527954  
Loss from our model:          0.08793757110834122
```

Standard deviation difference

```
In [18]: # switch back to cpu for plotting  
X_train = X_train.cpu()  
y_train = y_train.cpu()  
X_val = X_val.cpu()  
y_val = y_val.cpu()  
pred = pred.cpu()  
  
# backprop components no longer needed  
X_train = X_train.detach()  
y_train = y_train.detach()  
X_val = X_val.detach()  
y_val = y_val.detach()  
pred = pred.detach()
```

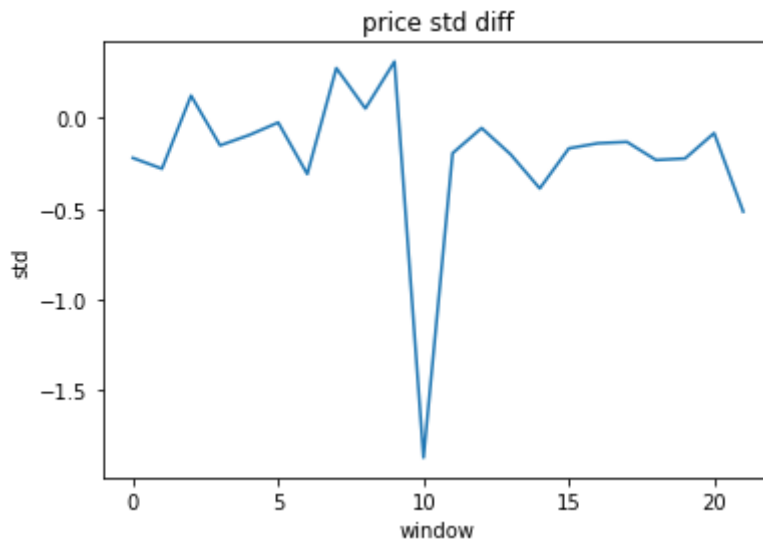


```
In [19]: f1 = plt.figure()

ax1 = f1.add_subplot()
ax1.plot((pred - y_val)[500:522])
ax1.set_title('price std diff')
ax1.set(xlabel='window', ylabel='std')

plt.show()

# plt.plot((pred[:,3] - y_val.cpu()[:,3]).detach())
# plt.title('std difference')
# plt.plot([1, 2, 3])
```



```
In [20]: # denormalize the data
pred_abs = pred * val_std[:,0].unsqueeze(1) + val_mean[:,0].unsqueeze(1)
y_val_abs = y_val.cpu() * val_std[:,0].unsqueeze(1) + val_mean[:,0].unsq
ueeze(1)
```

```
In [21]: start_window = 500

fig, (ax1, ax2) = plt.subplots(2, figsize=(15, 10))
ax1.plot((pred_abs - y_val_abs)[start_window:start_window + 60])
l1, = ax2.plot(pred_abs[start_window:start_window + 60])
l1.set_label("predicted price")
l2, = ax2.plot(y_val_abs[start_window:start_window + 60])
l2.set_label("actual price")

plt.legend()
plt.show()
```

