

RNN Moneymaker

```
In [1]: import torch
import torch.optim
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader

import time
import numpy as np
import matplotlib.pyplot as plt
```

Load data

```
In [2]: X_train, X_val, y_train, y_val, train_mean, val_mean, train_std, val_std
= torch.load("assets/all.pt")
```

```
In [3]: print("X_train shape: \t\t", X_train.shape)
print("X_val shape: \t\t", X_val.shape)
print("y_train shape: \t\t", y_train.shape)
print("y_val shape: \t\t", y_val.shape)
print("")

print("train_mean shape: \t", train_mean.shape)
print("val_mean: \t\t", val_mean.shape)
print("train_std: \t\t", train_std.shape)
print("val_std: \t\t", val_std.shape)
print("")
```

```
X_train shape:          torch.Size([2542795, 122, 3])
X_val shape:            torch.Size([282732, 122, 3])
y_train shape:          torch.Size([2542795, 1])
y_val shape:            torch.Size([282732, 1])

train_mean shape:       torch.Size([2542795, 3])
val_mean:               torch.Size([282732, 3])
train_std:               torch.Size([2542795, 3])
val_std:                 torch.Size([282732, 3])
```

```
In [4]: # optimize for GPU if exists
if torch.cuda.is_available():
    X_train = X_train.cuda()
    y_train = y_train.cuda()
    X_val = X_val.cuda()
    y_val = y_val.cuda()
```

Training a RNN

```
In [5]: class Dataset(torch.utils.data.Dataset):
        def __init__(self, X, y):
            self.X = X
            self.y = y

        def __len__(self):
            return len(self.X)

        def __getitem__(self, index):
            return self.X[index], self.y[index]
```

```
In [6]: batch_size = 32
        dataset = Dataset(X_train, y_train)
        loader = DataLoader(dataset, batch_size, shuffle=True)
```

```
In [7]: class RNNClassifier(nn.Module):
        def __init__(self, input_dim, hidden_dim, output_dim):
            super(RNNClassifier, self).__init__()
            self.input_dim = input_dim
            self.hidden_dim = hidden_dim
            self.output_dim = output_dim

            self.rnn = nn.LSTM(input_dim, hidden_dim, batch_first=True, num_
layers=3)
            self.drop = nn.Dropout(0.5)
            self.fc = nn.Linear(hidden_dim, output_dim)

        def forward(self, x, h=None):
            if type(h) == type(None):
                out, hn = self.rnn(x)
            else:
                out, hn = self.rnn(x, h.detach())
            out = self.drop(out)
            out = self.fc(out[:, -1, :])
            return out
```

```
In [8]: input_dim = 3
        hidden_dim = 20
        output_dim = 1
```

```
In [9]: model = RNNClassifier(input_dim, hidden_dim, output_dim)
        if torch.cuda.is_available():
            model = model.to("cuda")

        criterion = nn.MSELoss()
        optimizer = torch.optim.Adam(model.parameters())
```

```
In [10]: train_accs = []
         val_accs = []
         train_losses = []
         val_losses = []
         epoch = 0
```

Train the model

```
In [11]: def pred_val(X_val, model):
         val_batch_size = 1000
         val_set_size = X_val.shape[0]
         preds = []
         with torch.no_grad():
             for i in range(0, val_set_size, val_batch_size):
                 start = i
                 end = min(i+val_batch_size, val_set_size)
                 preds.append(model(X_val[start:end]))
         pred = torch.cat(preds, dim=0)
         return pred
```

```

In [12]: t0 = time.time()
num_epochs = 3
for ep in range(num_epochs):
    tstart = time.time()
    for i, data in enumerate(loader):
        model.train()
        print("{} / {}".format(i, len(loader)), end='\r')
        optimizer.zero_grad()
        outputs = model(data[0])
        loss = criterion(outputs, data[1])
        loss.backward()
        optimizer.step()

    if i % 1000 == 0:
        with torch.no_grad():
            model.eval()
            train_losses.append(loss.item())
            pXval = pred_val(X_val, model)
            vloss = criterion(pXval, y_val)
            val_losses.append(vloss.item())
            torch.save({
                'epoch': epoch,
                'model_state_dict': model.state_dict(),
                'optimizer_state_dict': optimizer.state_dict(),
                'loss': loss,
            }, 'assets/partial_model.pt')

        print("training loss: {:<3.3f} \t val loss: {:<3.3f}".fo
rmat(loss, vloss))

    with torch.no_grad():
        model.eval()
        pXval = model(X_val)
        vloss = criterion(pXval, y_val)
        val_losses.append(vloss.item())
        epoch += 1
        tend = time.time()
        print('epoch: {:<3d} \t time: {:<3.2f} \t val loss: {:<3.3f}'.fo
rmat(epoch,
        tend - tstart, vloss.item()))
time_total = time.time() - t0
print('Total time: {:<4.3f}, average time per epoch: {:<4.3f}'.format(time
_total, time_total / num_epochs))

```

training loss: 1.701	val loss: 2.028
training loss: 0.168	val loss: 0.163
training loss: 0.137	val loss: 0.153
training loss: 0.117	val loss: 0.164
training loss: 0.181	val loss: 0.155
training loss: 0.148	val loss: 0.154
training loss: 0.167	val loss: 0.149
training loss: 0.157	val loss: 0.146
training loss: 0.121	val loss: 0.147
training loss: 0.138	val loss: 0.146
training loss: 0.115	val loss: 0.152
training loss: 0.511	val loss: 0.161
training loss: 0.202	val loss: 0.149
training loss: 0.185	val loss: 0.170
training loss: 0.175	val loss: 0.149
training loss: 0.317	val loss: 0.146
training loss: 0.083	val loss: 0.145
training loss: 0.630	val loss: 0.143
training loss: 0.178	val loss: 0.143
training loss: 0.109	val loss: 0.146
training loss: 0.108	val loss: 0.145
training loss: 0.317	val loss: 0.145
training loss: 0.352	val loss: 0.147
training loss: 0.173	val loss: 0.147
training loss: 0.222	val loss: 0.154
training loss: 0.421	val loss: 0.149
training loss: 0.374	val loss: 0.144

26541/79463

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-12-26b9ada45317> in <module>
      7         print("{} / {}".format(i, len(loader)), end='\r')
      8         optimizer.zero_grad()
---->  9         outputs = model(data[0])
     10         loss = criterion(outputs, data[1])
     11         loss.backward()

/opt/anaconda3/lib/python3.7/site-packages/torch/nn/modules/module.py in
n __call__(self, *input, **kwargs)
     530         result = self._slow_forward(*input, **kwargs)
     531     else:
-->  532         result = self.forward(*input, **kwargs)
     533     for hook in self._forward_hooks.values():
     534         hook_result = hook(self, input, result)

<ipython-input-7-2e69c636e5f9> in forward(self, x, h)
     12     def forward(self, x, h=None):
     13         if type(h) == type(None):
---->  14             out, hn = self.rnn(x)
     15     else:
     16         out, hn = self.rnn(x, h.detach())

/opt/anaconda3/lib/python3.7/site-packages/torch/nn/modules/module.py in
n __call__(self, *input, **kwargs)
     530         result = self._slow_forward(*input, **kwargs)
     531     else:
-->  532         result = self.forward(*input, **kwargs)
     533     for hook in self._forward_hooks.values():
     534         hook_result = hook(self, input, result)

/opt/anaconda3/lib/python3.7/site-packages/torch/nn/modules/rnn.py in f
orward(self, input, hx)
     557         if batch_sizes is None:
     558             result = _VF.lstm(input, hx, self._flat_weights, se
lf.bias, self.num_layers,
-->  559                                     self.dropout, self.training, sel
f.bidirectional, self.batch_first)
     560     else:
     561         result = _VF.lstm(input, batch_sizes, hx, self._fla
t_weights, self.bias,


```

KeyboardInterrupt:

Training loss vs. validation loss

```
In [13]: t_losses = [i for i in train_losses if i < 4000]
plt.plot(t_losses)
plt.plot(val_losses)
plt.title('loss history')
plt.xlabel('thousand iterations')
plt.ylabel('loss')
plt.yscale('log')
plt.legend(['train', 'val'])

plt.show()
```



Evaluate the model

```
In [14]: model.eval()

pred = pred_val(X_val, model)
val_loss = criterion(pred, y_val).item()

print("Final model evaluation: ", val_loss)
```

Final model evaluation: 0.15096266567707062

One-step lag predictor

The one-step lag predictor simply outputs the last timestep in the input sequence. Our model should outperform the one-step lag predictor.

```
In [15]: def one_step_lag_predictor(X):  
         return X[:, -1, 0].unsqueeze(1)  
  
p_val_naive = one_step_lag_predictor(X_val.cpu())  
loss_naive = criterion(p_val_naive, y_val.cpu())  
  
print("Loss from 1-step lag predictor:\t{}\nLoss from our model:\t\t{}".  
      format(loss_naive, val_loss))
```

```
Loss from 1-step lag predictor: 0.14193207025527954  
Loss from our model:          0.15096266567707062
```

Standard deviation difference

```
In [16]: # switch back to cpu for plotting  
X_train = X_train.cpu()  
y_train = y_train.cpu()  
X_val = X_val.cpu()  
y_val = y_val.cpu()  
pred = pred.cpu()  
  
# backprop components no longer needed  
X_train = X_train.detach()  
y_train = y_train.detach()  
X_val = X_val.detach()  
y_val = y_val.detach()  
pred = pred.detach()
```




```
In [17]: f1 = plt.figure()

ax1 = f1.add_subplot()
ax1.plot((pred - y_val)[500:522])
ax1.set_title('price std diff')
ax1.set(xlabel='window', ylabel='std')

plt.show()

# plt.plot((pred[:,3] - y_val.cpu()[:,3]).detach())
# plt.title('std difference')
# plt.plot([1, 2, 3])
```



Actual price difference

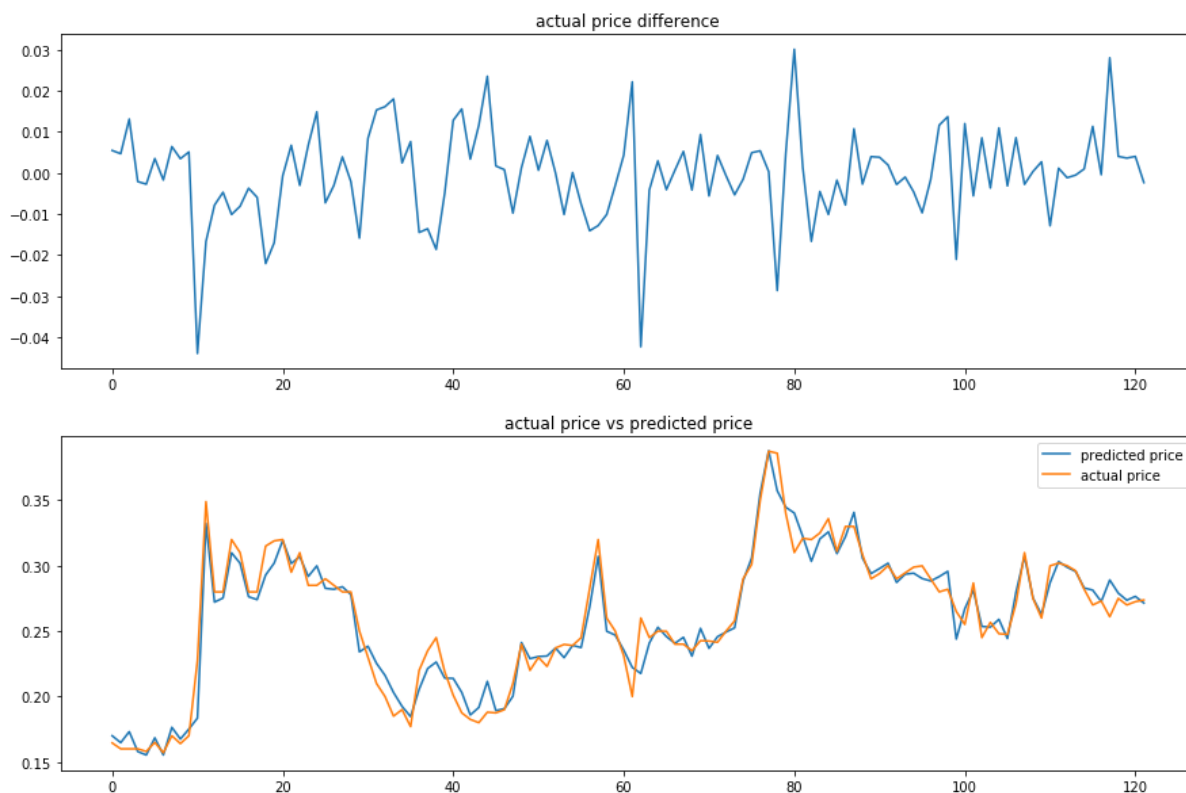
```
In [18]: # denormalize the data
pred_abs = pred * val_std[:,0].unsqueeze(1) + val_mean[:,0].unsqueeze(1)
y_val_abs = y_val.cpu() * val_std[:,0].unsqueeze(1) + val_mean[:,0].unsqueeze(1)
```

```
In [19]: fig, (ax1, ax2) = plt.subplots(2, figsize=(15, 10))

ax1.set_title("actual price difference")
ax1.plot((pred_abs - y_val_abs)[500:622])

ax2.set_title("actual price vs predicted price")
l1, = ax2.plot(pred_abs[500:622])
l1.set_label("predicted price")
l2, = ax2.plot(y_val_abs[500:622])
l2.set_label("actual price")

plt.legend()
plt.show()
```



```
In [23]: fig, ax = plt.subplots(1, figsize=(8, 5))
ax.set_title("actual price vs predicted price")
l1, = ax.plot(pred_abs[300:380])
l1.set_label("predicted price")
l2, = ax.plot(y_val_abs[300:380])
l2.set_label("actual price")

plt.legend()
plt.show()
```

